

PertAdapt: Unlocking Single-Cell Foundation Models for Genetic Perturbation Prediction via Condition-Sensitive Adaptation

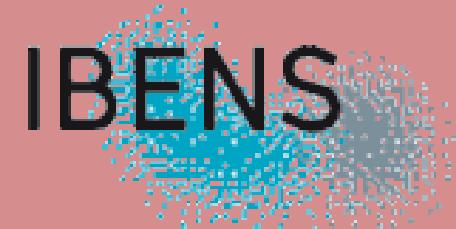
Ding Bai, Le Song, Eric Xing

Published in



(November 21st, 2025)

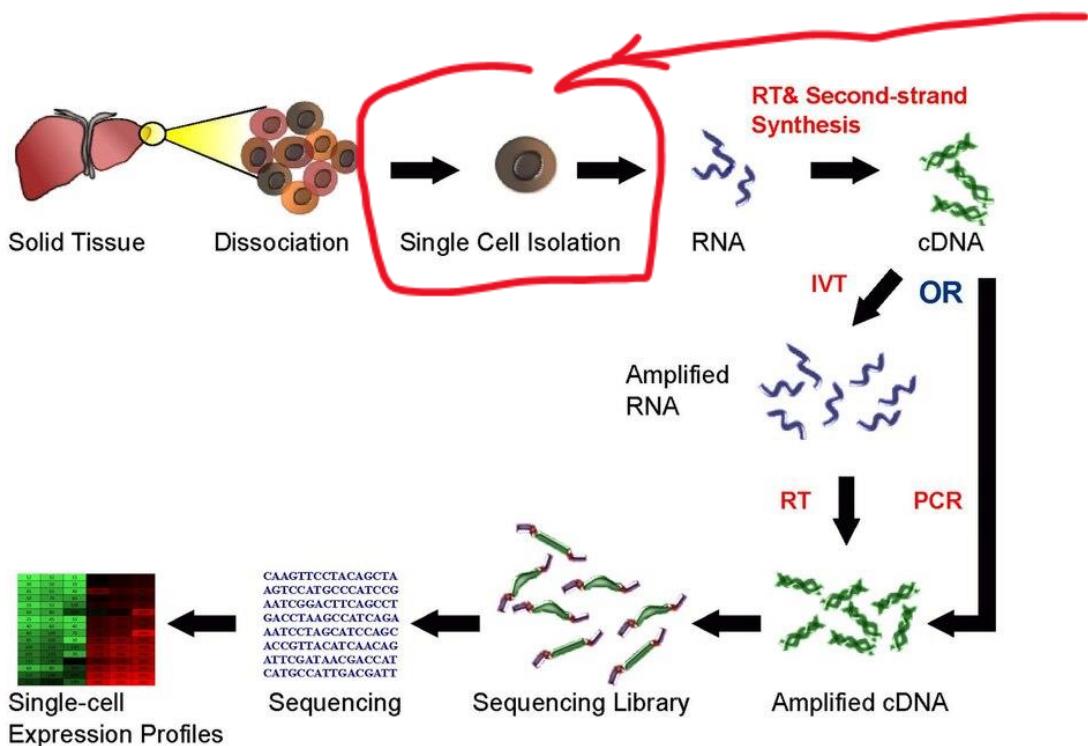
03.03.2026 Clemence Reda



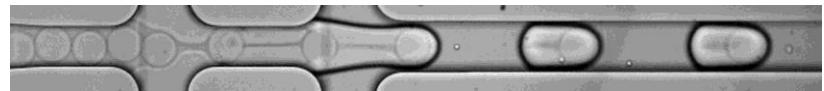
Background

1. Single-cell RNA sequencing (very short intro)
2. Gene Ontology annotations
3. Graph (Convolutional) Neural Networks (very short intro)
4. Problem setting
5. SOTA model GEARS (actually not so SOTA...)

Single-cell RNA sequencing snapshots of cell-specific expression in a tissue



Source from Wikipedia ("Single-cell sequencing")



Specific to single-cell seq
(e.g., droplet)

GEM-X from 10x Genomics [1]
500-20k cells/tissue
~2-4 days end-to-end
8-10k per-cell gene counts
~\$3k-\$8k per sample [2]

- [1] 10xgenomics.com/blog/the-next-generation-of-single-cell-rna-seq-an-introduction-to-gem-x-technology (May, 2024)
[2] satijalab.org/costpercell/

Gene Ontology (GO) annotations [3] Contribution of individual gene (products) to biological functions

One gene can be attached to multiple GO annotations

GO annotation

- Gene product
- GO term
- Reference paper
- GO Evidence code

Biological Process: e.g., DNA repair

Molecular Function: e.g., catalysis

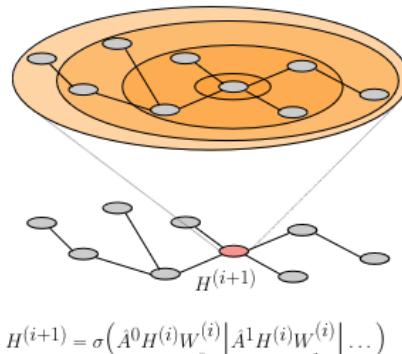
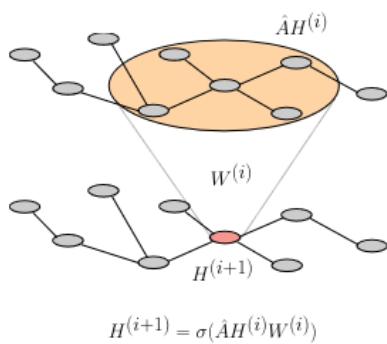
Cellular Component: e.g., cytoskeleton

Existence of a **hierarchy** for GO terms

What people do: count the number of shared GO terms in a pair of genes

[3] The Gene Ontology Consortium. (2025) *Nucleic Acids Research*;gkaf1292.

Graph (Convolutional) Neural Networks [4] Get node representations infused with graph information



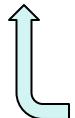
$H^{(i)}$: embedding matrix in $\mathbb{R}^{N \times d}$ at neural network layer i

in practice $H^{(i+1)} = \hat{A}^k H^{(i)} W^{(i)}$ *& normalized adj. matrix*

K=1 versus K>1: adapted from [5]

Spatial domain $H^{(i)}_v = f(\text{Agg}(\{H^{(i-1)}_u \mid (u,v) \text{ in } G\})W^{(i)}, H^{(i-1)}_v)$ Choice of the aggregation function? [6]

Spectral domain $H^{(i)} = U H^{(i-1)} W^{(i)}$



Spectral filter: eigenvectors of Laplacian of graph
 $G = \text{Degree}(G) - \text{Adjacency}(G)$

U can be numerically instable... [7]

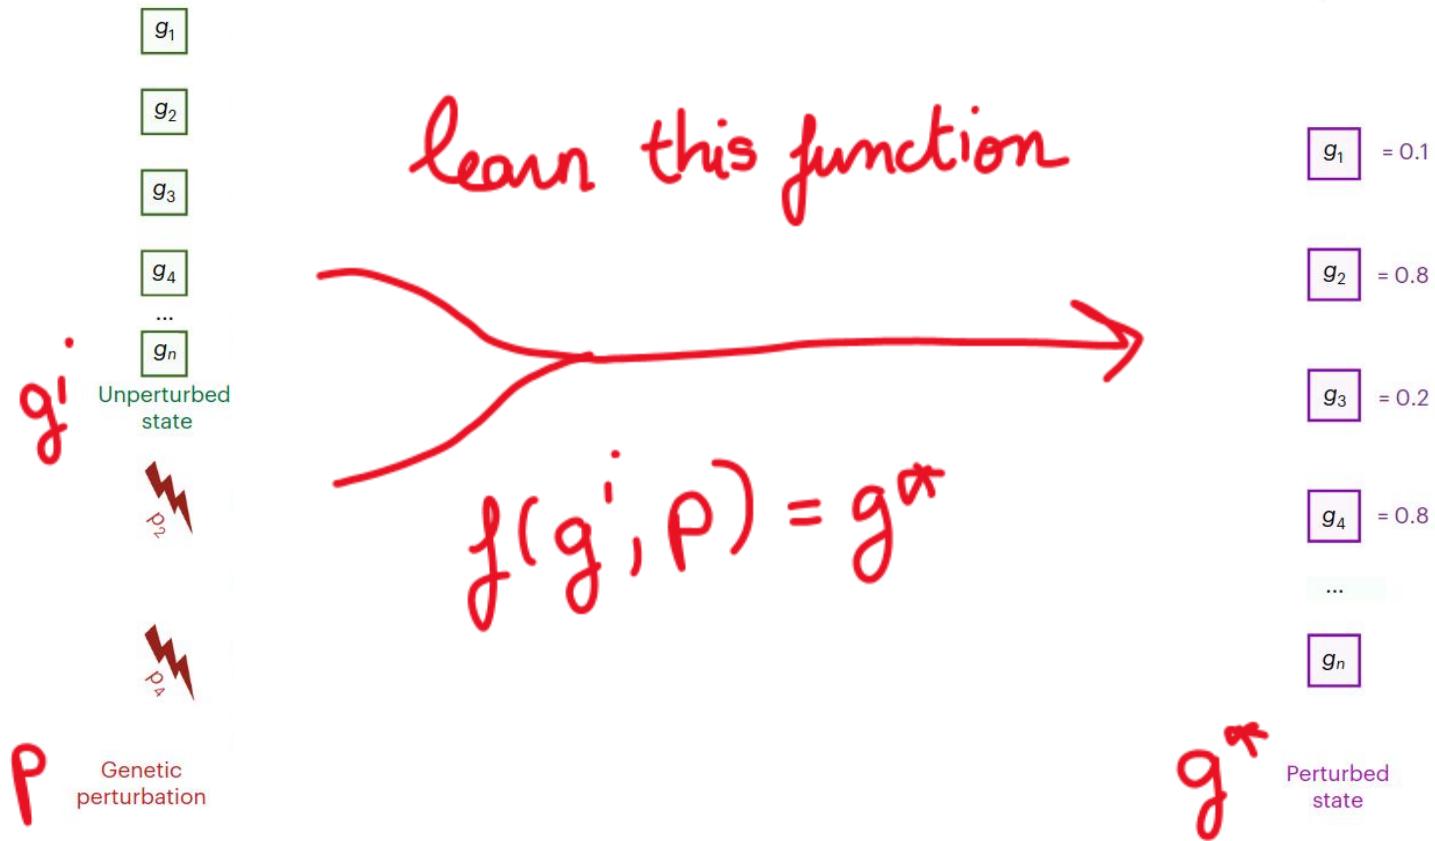
[4] Kipf and Welling. (2016) *International Conference on Learning Representations*

[5] Abu-El-Haija et al. (2019) *International Conference on Machine Learning*

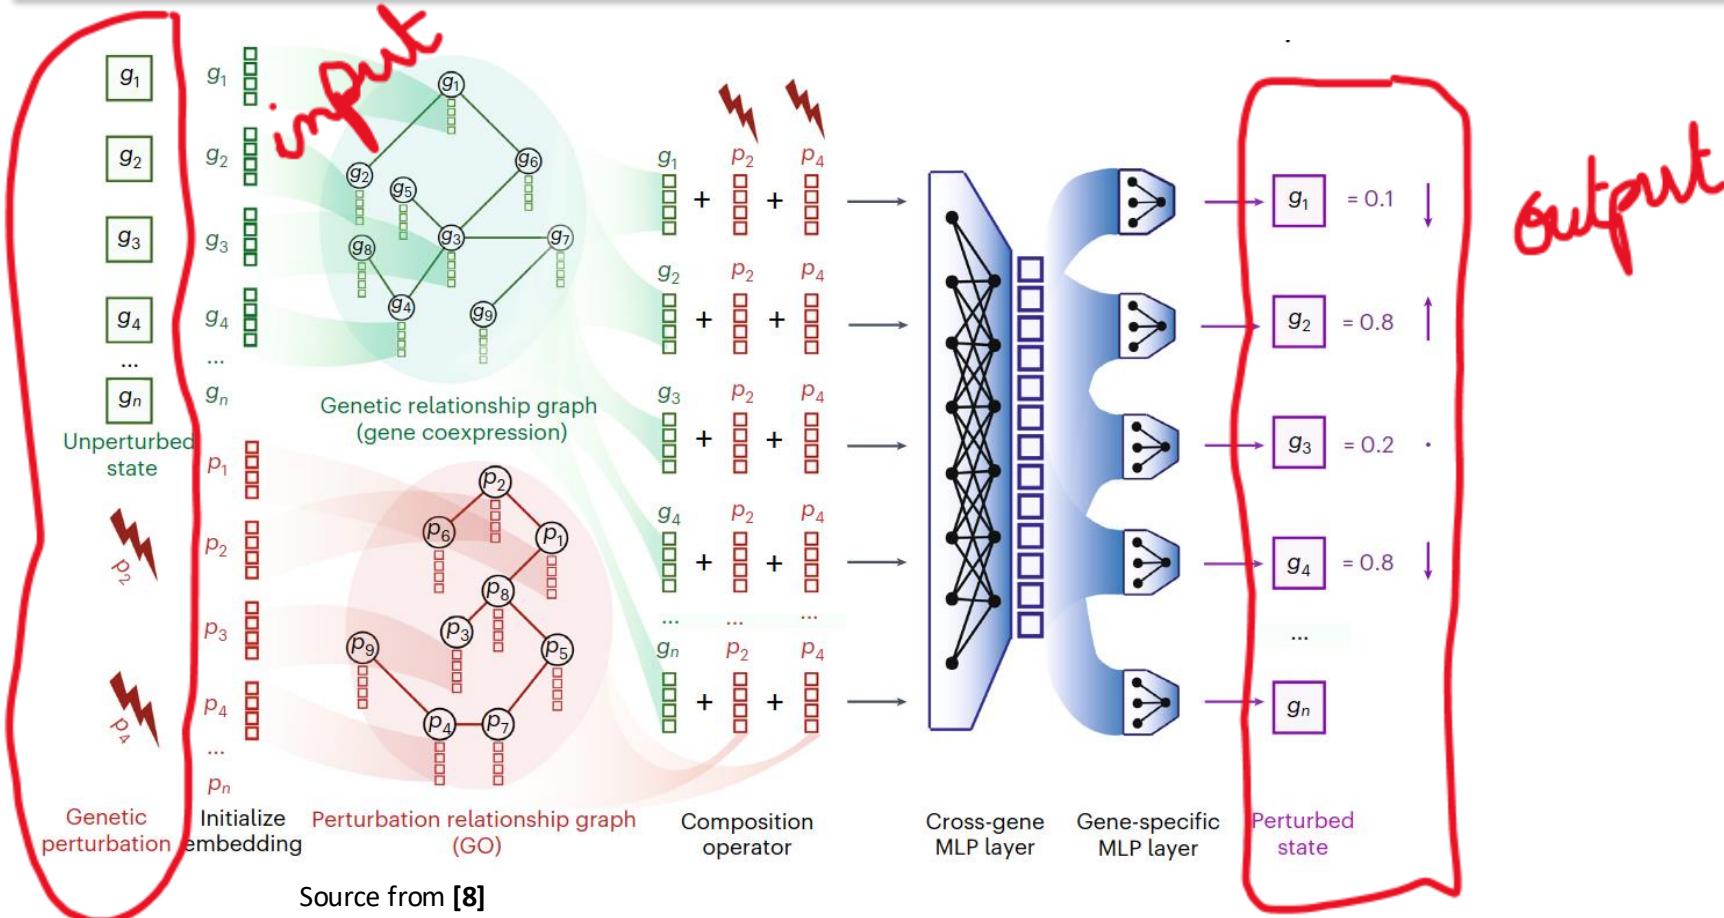
[6] Xu et al. (2019) *International Conference on Learning Representations*

[7] Park et al. (2019) *IEEE/CVF Conference on Computer Vision and Pattern Recognition*

Problem setting Prediction of post-perturbation (on >1 gene) single-cell expression profile

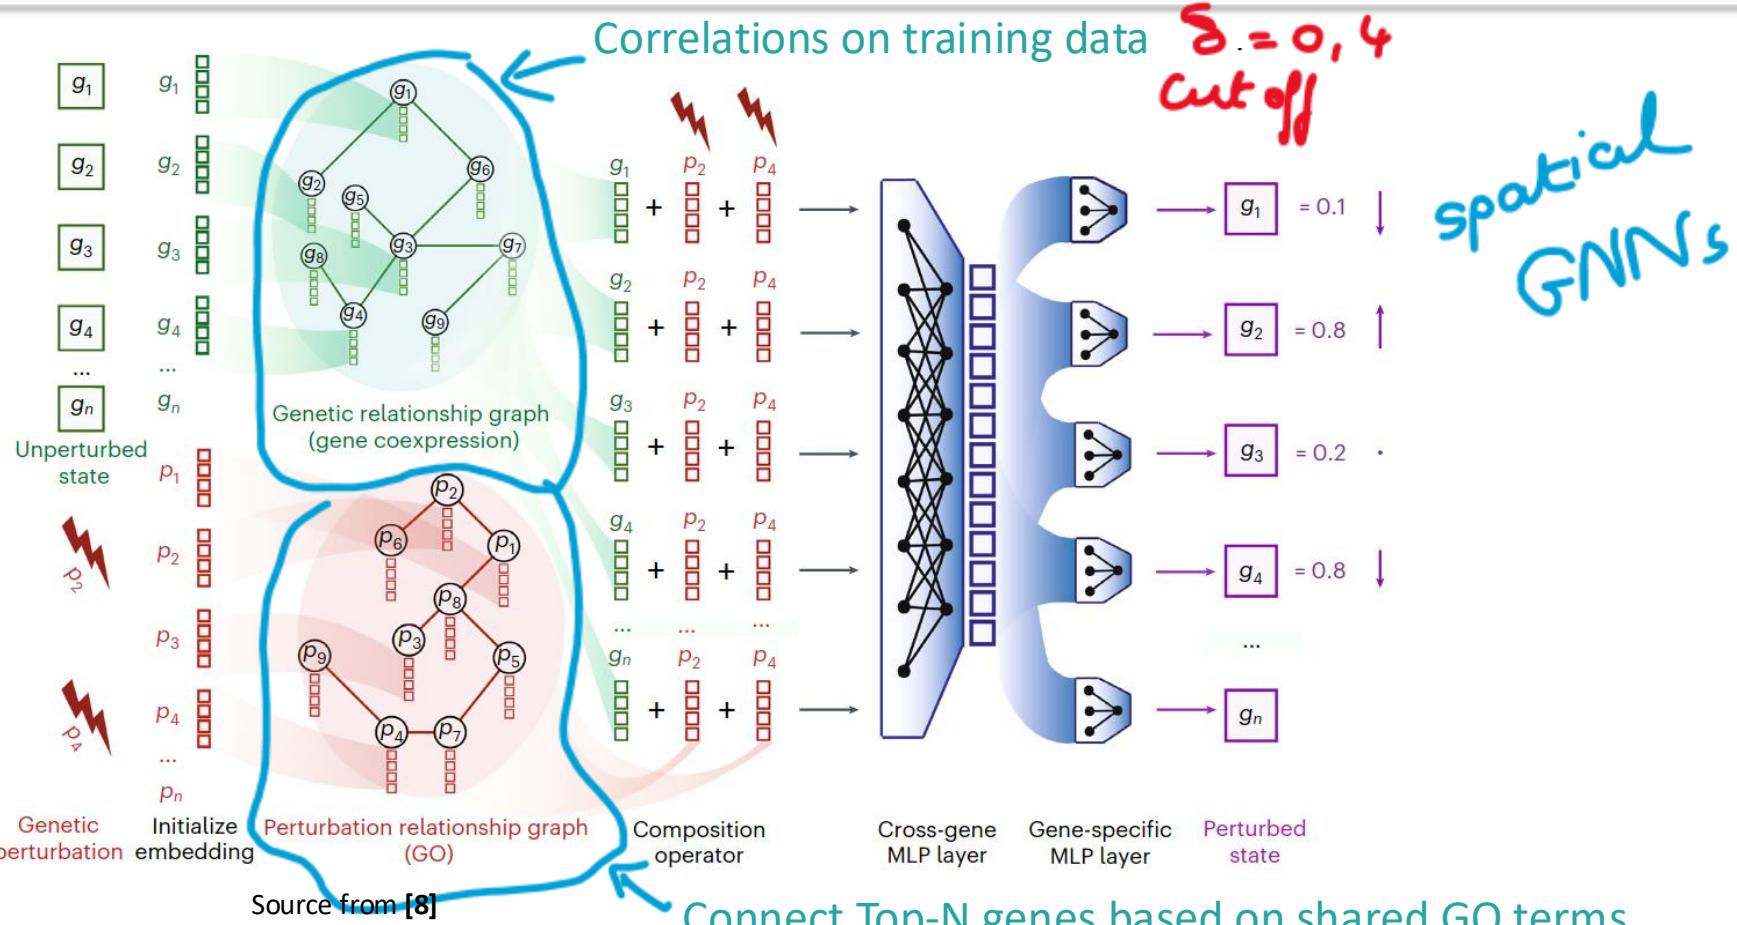


SOTA model GEARS [8] predicting post-perturbation single-cell expression profiles

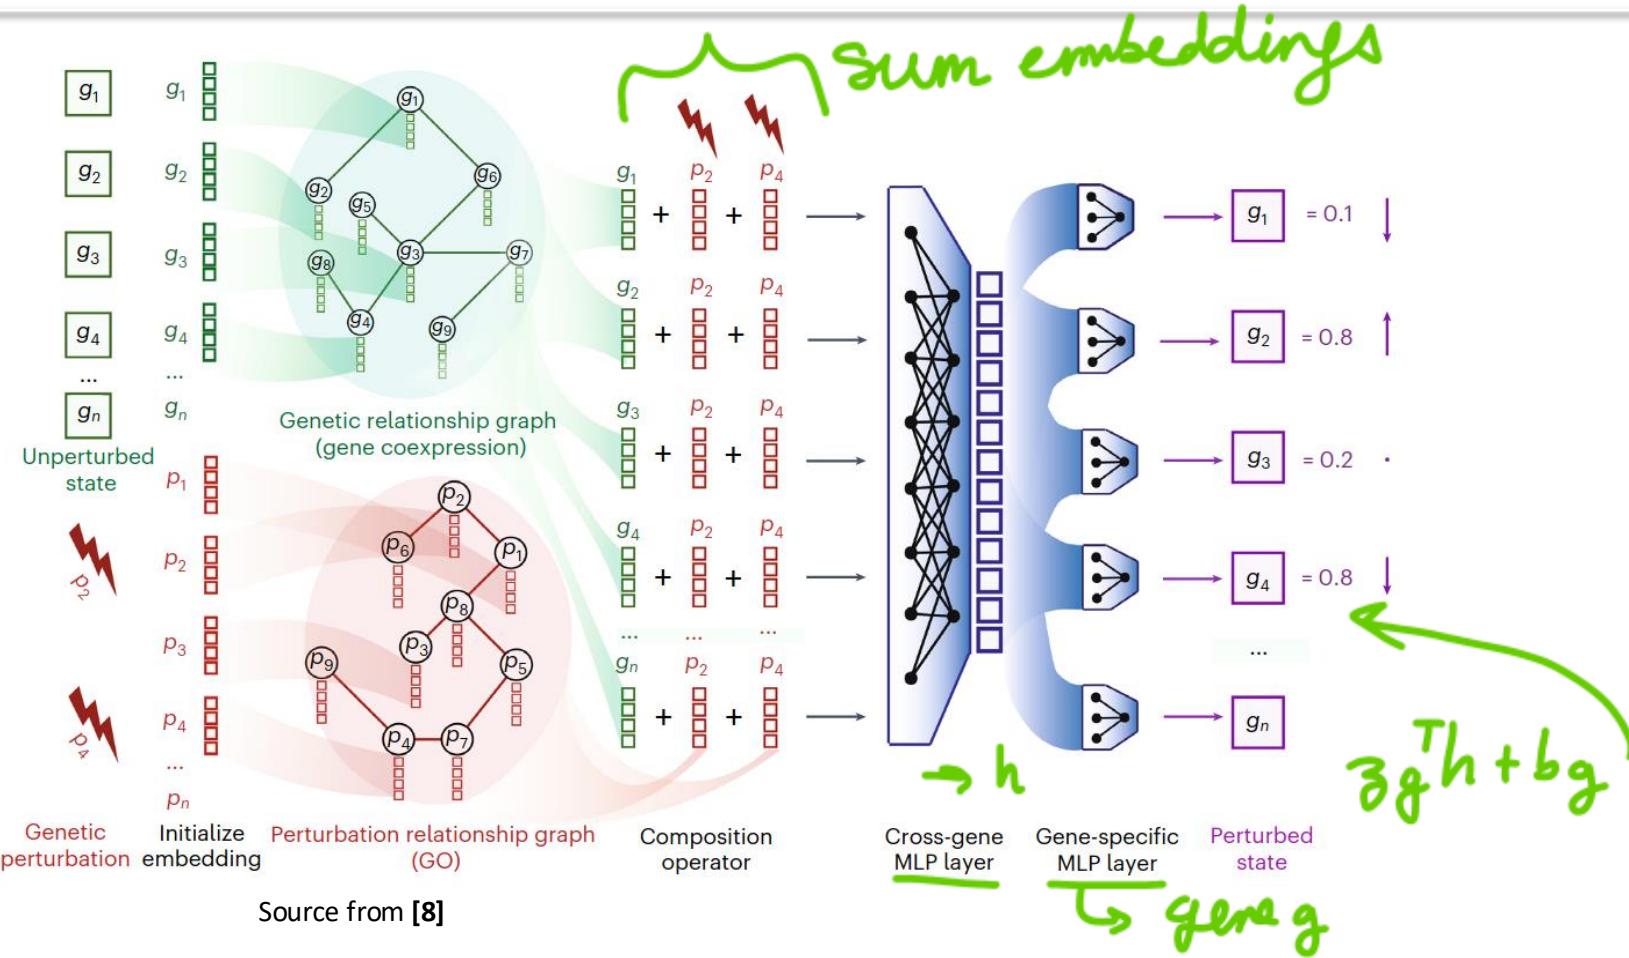


[8] Roohani, Huang & Leskovec. (2023) *Nature Biotechnology* (1):9.

SOTA model GEARS [8] predicting post-perturbation single-cell expression profiles

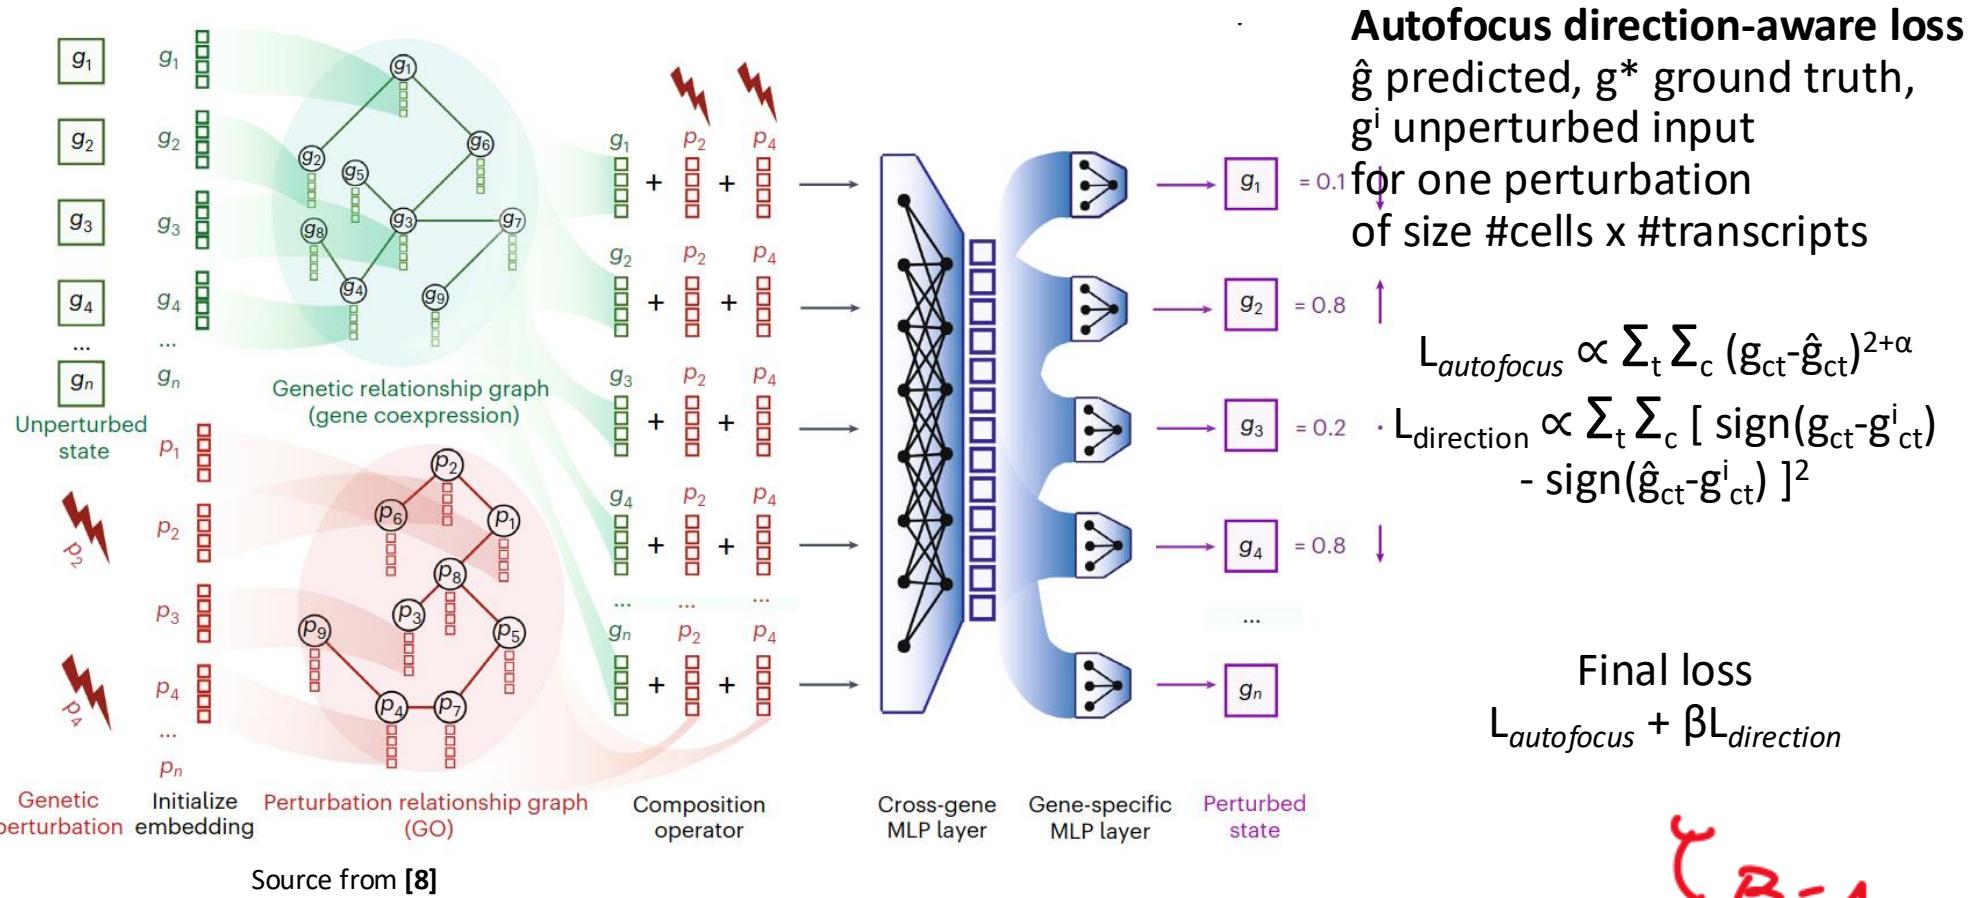


SOTA model GEARS [8] predicting post-perturbation single-cell expression profiles



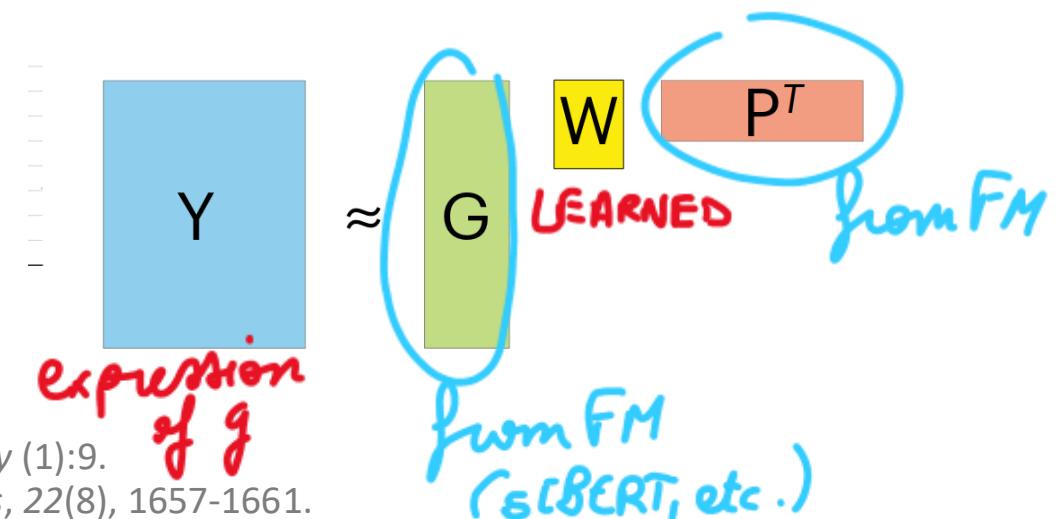
[8] Roohani, Huang & Leskovec. (2023) *Nature Biotechnology* (1):9.

SOTA model GEARS [8] predicting post-perturbation single-cell expression profiles



[8] Roohani, Huang & Leskovec. (2023) *Nature Biotechnology* (1):9.

SOTA model GEARS [8] Is it really state-of-the-art?



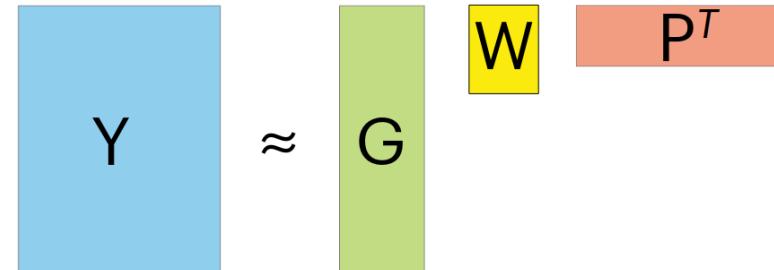
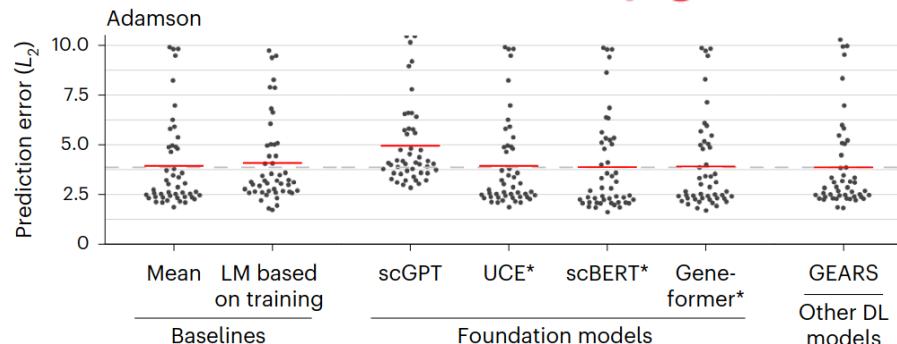
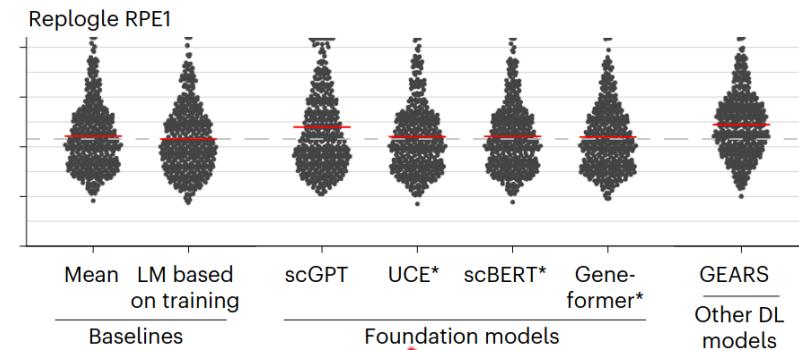
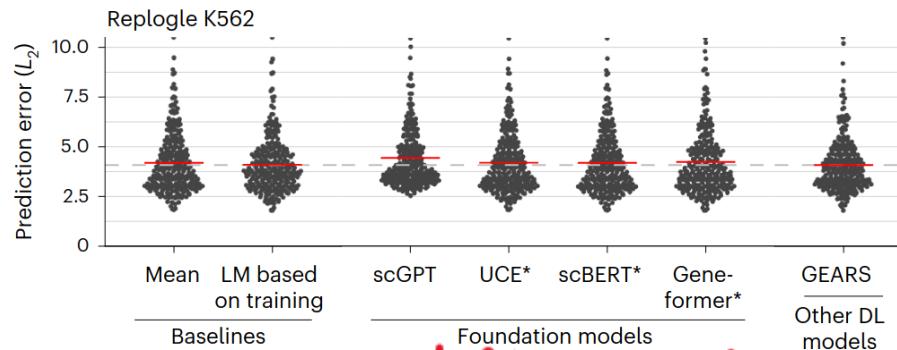
Source from [9]

[8] Roohani, Huang & Leskovec. (2023) *Nature Biotechnology* (1):9.

[9] Ahlmann-Eltze, Huber & Anders. (2025). *Nature Methods*, 22(8), 1657-1661.

SOTA model GEARS [8] Is it really state-of-the-art?

Does not consistently outperform linear regression... [9]



[8] Roohani, Huang & Leskovec. (2023) *Nature Biotechnology* (1):9.

[9] Ahlmann-Eltze, Huber & Anders. (2025). *Nature Methods*, 22(8), 1657-1661.

Content of the paper

"[We] propose PertAdapt, an FM-based method [...] that integrates a plug-in perturbation adapter with an adaptive loss"

Specific issues:

- Enabling downstream task-specific tuning instead of plug-in MLPs
- Reweighting gene importance according to differential expression

Prior approaches:

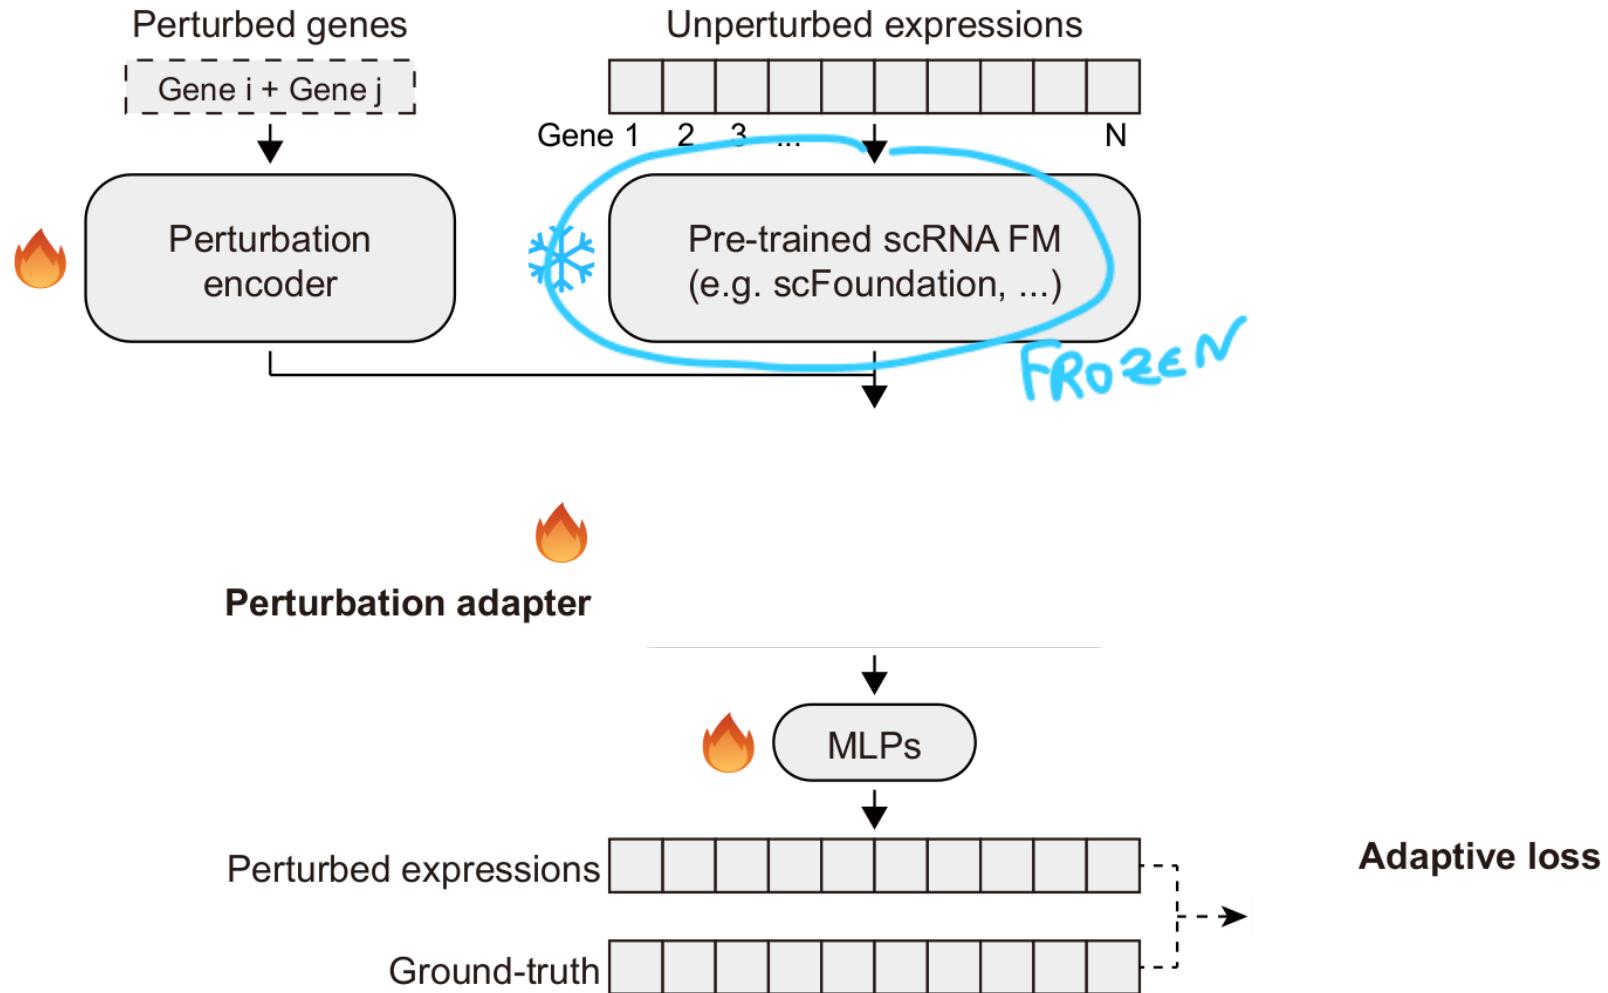
- SOTA GEARS, (repurposed) FMs: Geneformer, UCE, scFoundation,...

Content of the paper

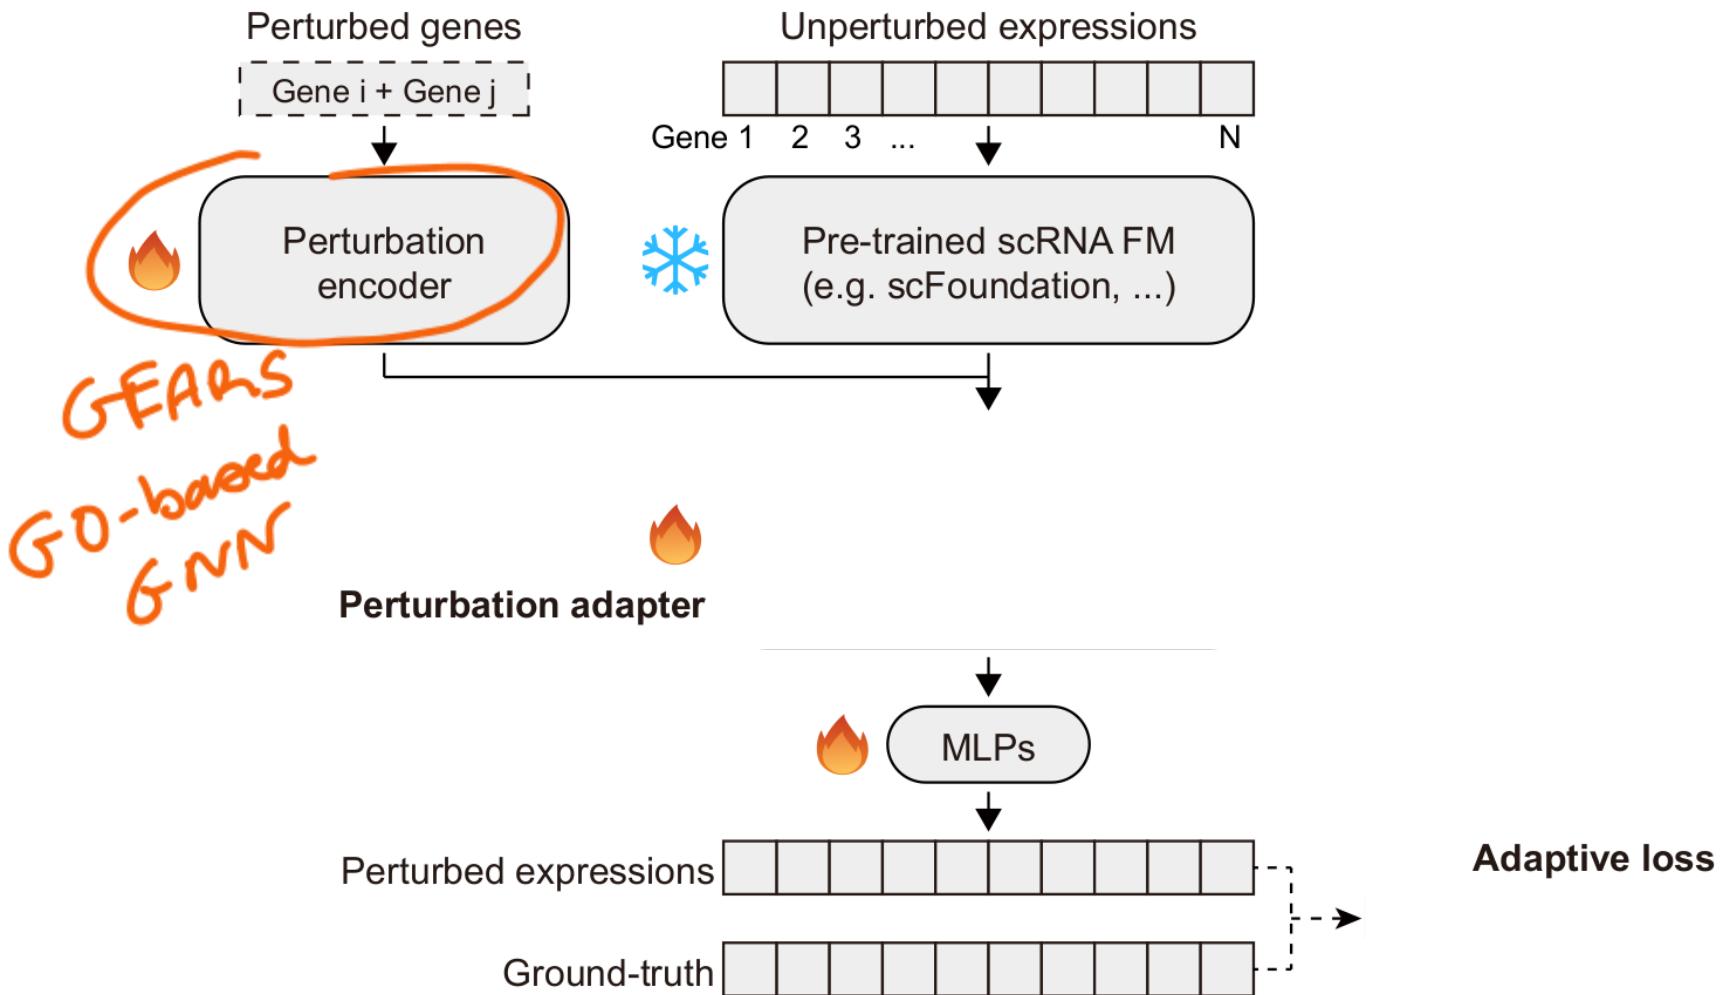
"[We] propose PertAdapt, an FM-based method [...] that integrates a plug-in perturbation adapter with an adaptive loss"

1. Architecture of PertAdapt
2. Plug-in perturbation adapter
3. Adaptive loss
4. Experiments: Prediction of post-perturbation expr.
5. Experiments: Ablation study

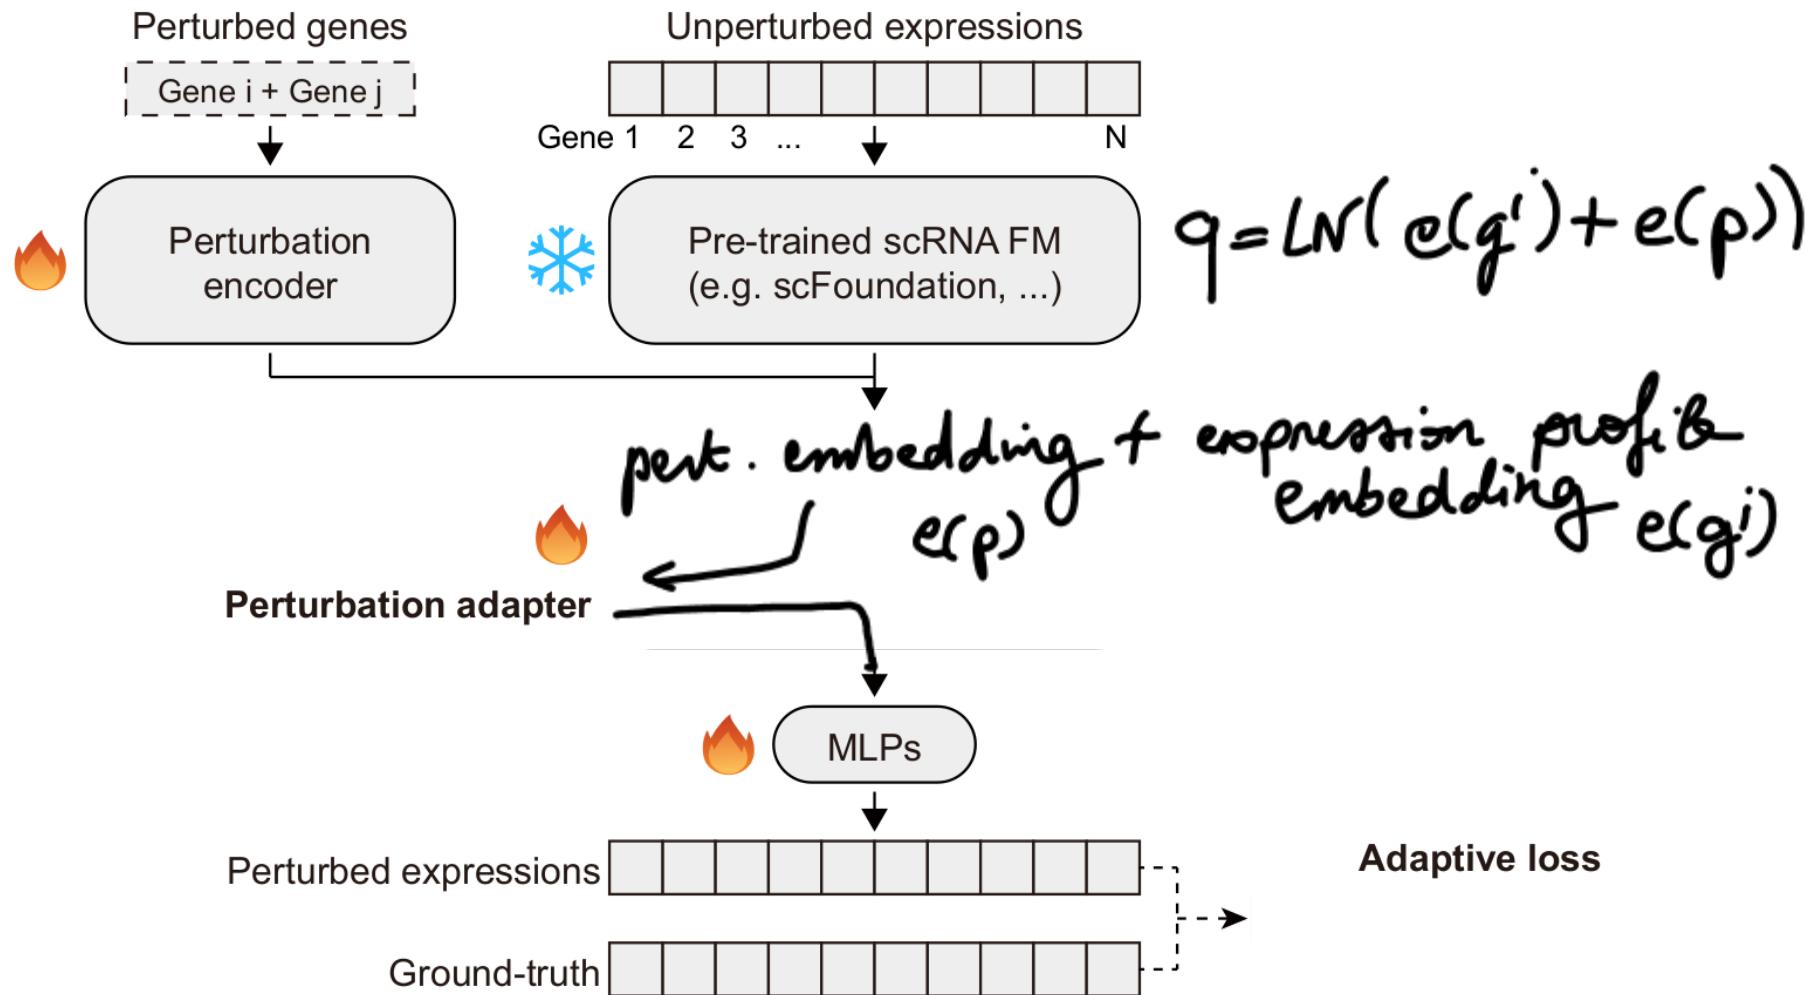
Architecture of PertAdapt foundation model-agnostic structure



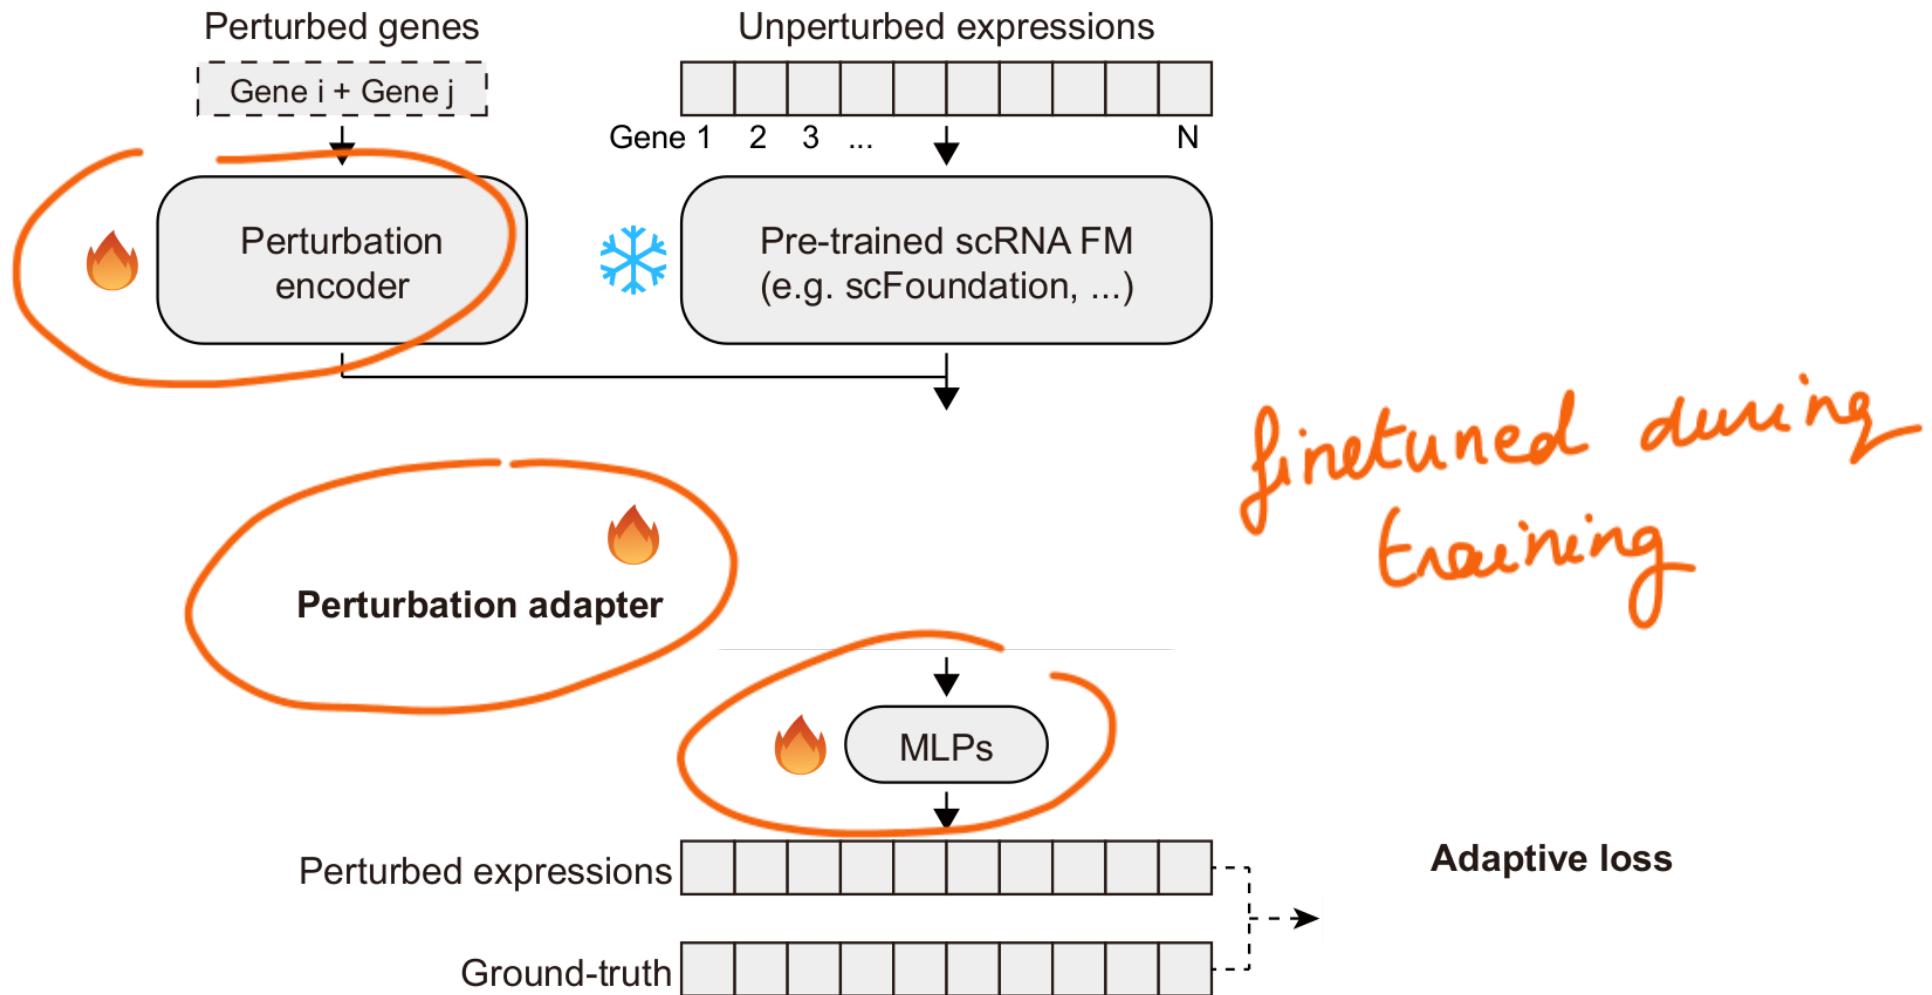
Architecture of PertAdapt foundation model-agnostic structure



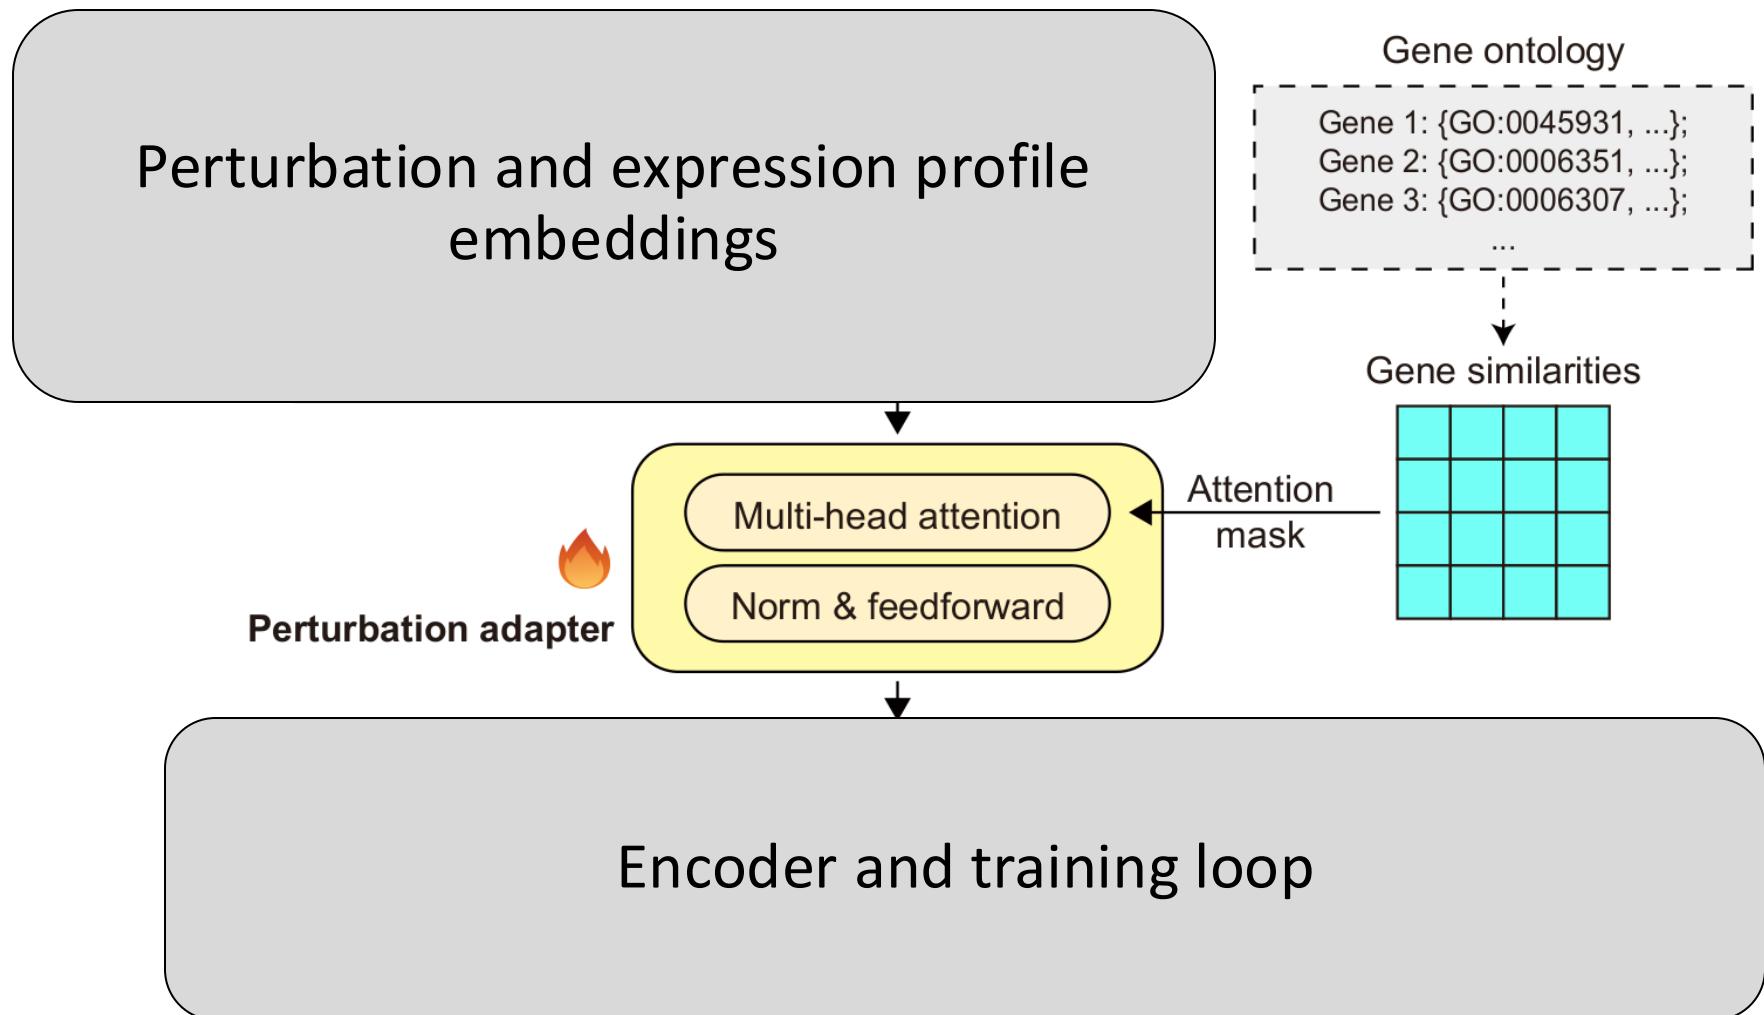
Architecture of PertAdapt foundation model-agnostic structure



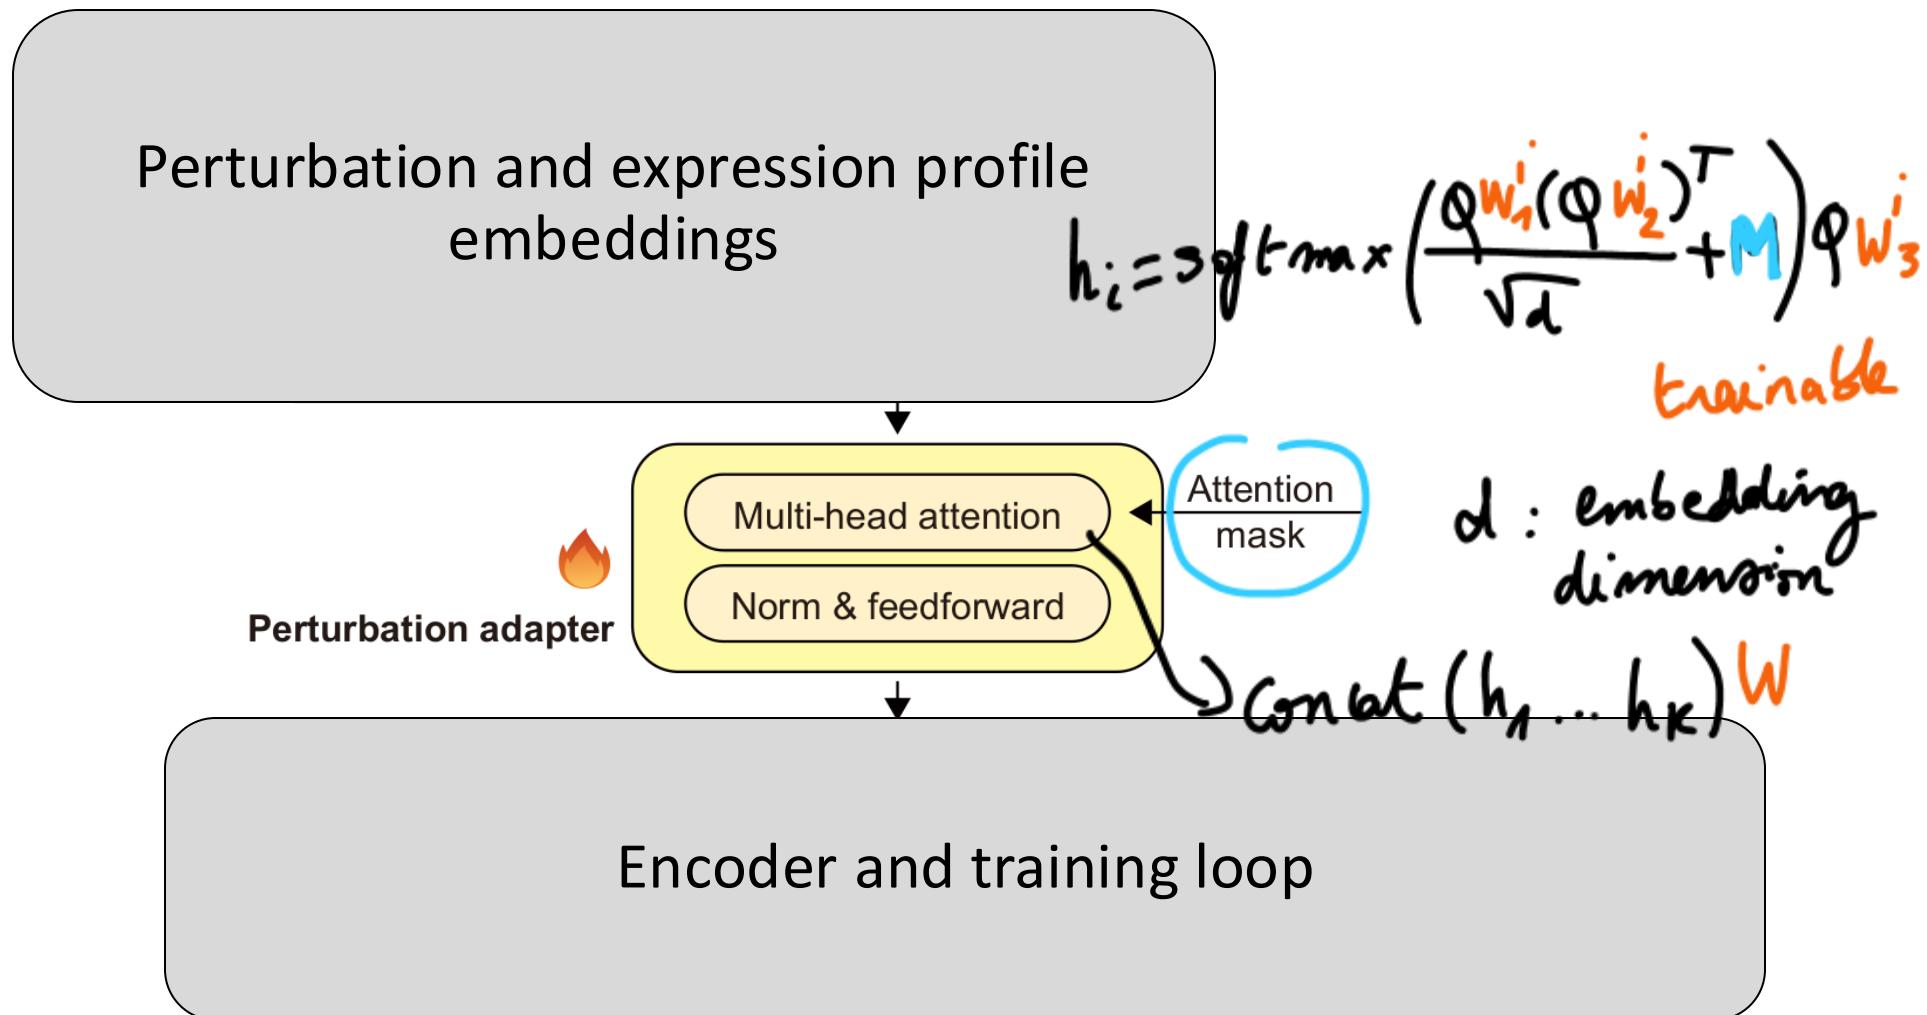
Architecture of PertAdapt foundation model-agnostic structure



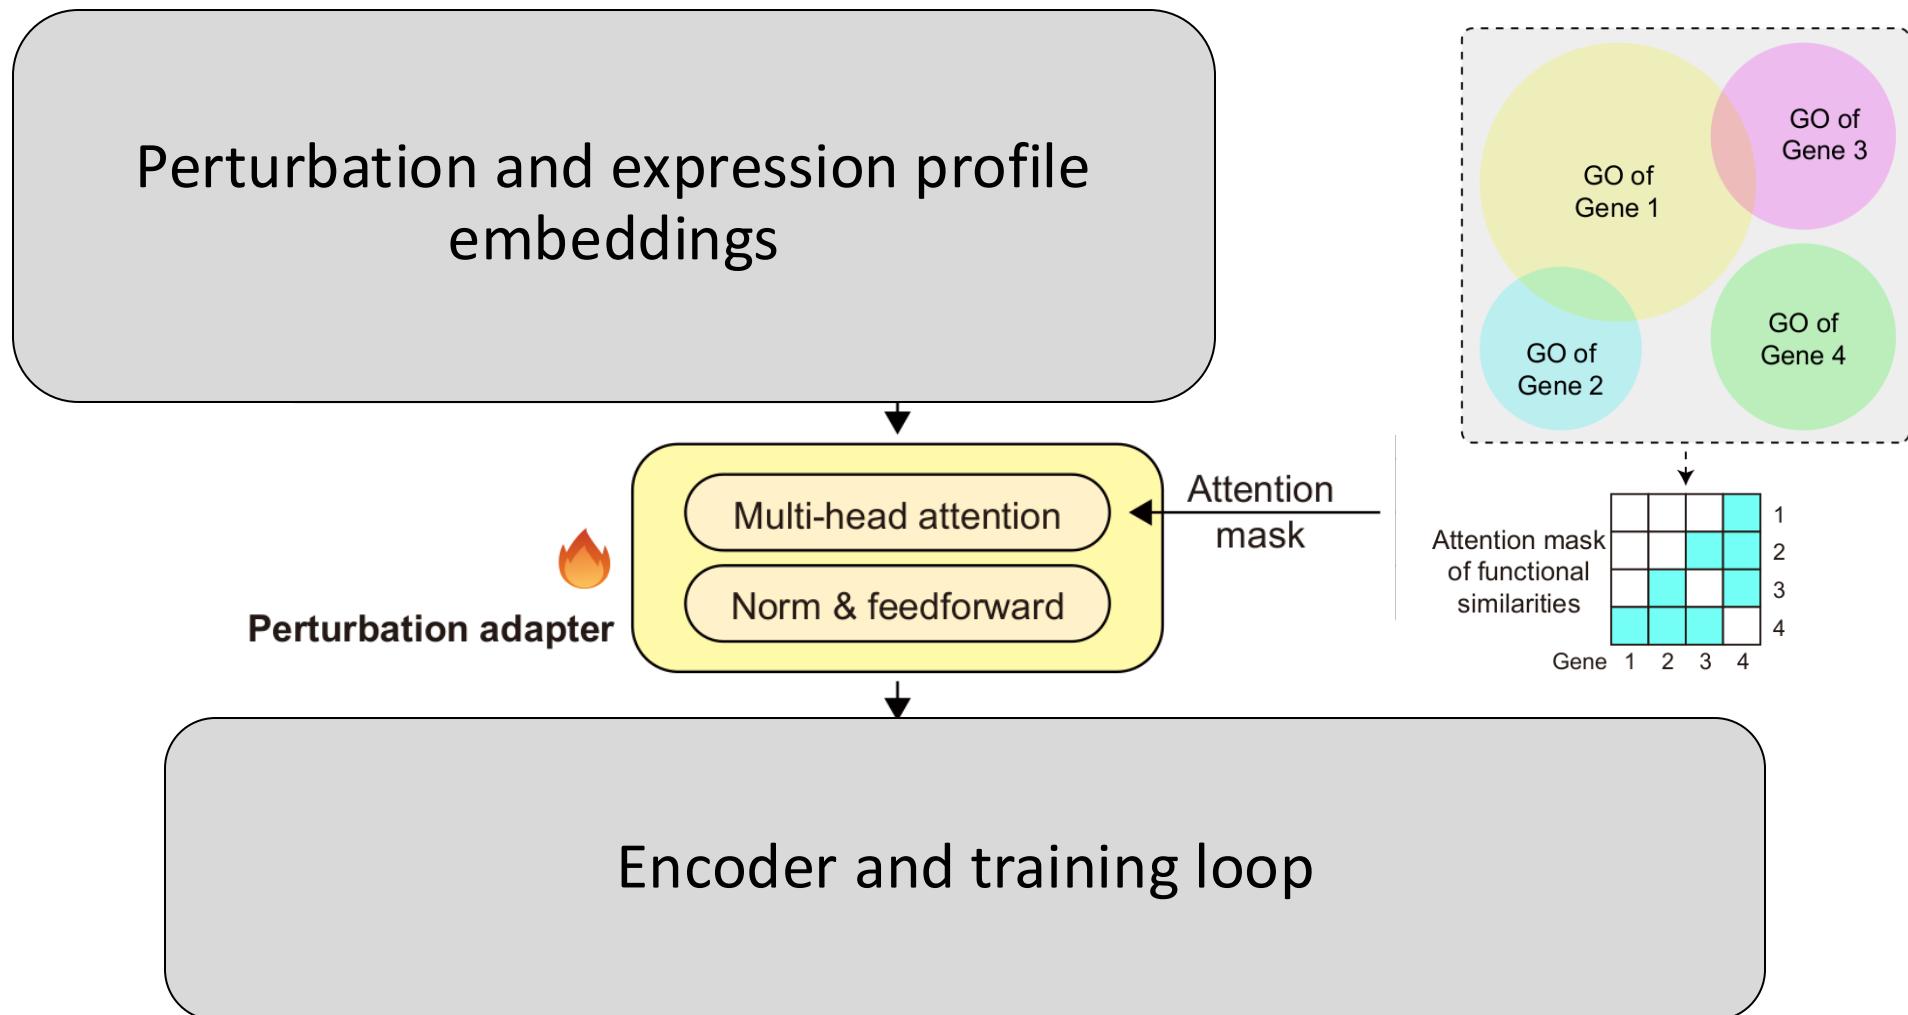
Plug-in perturbation adapter downstream-task tuning and biological prior integration



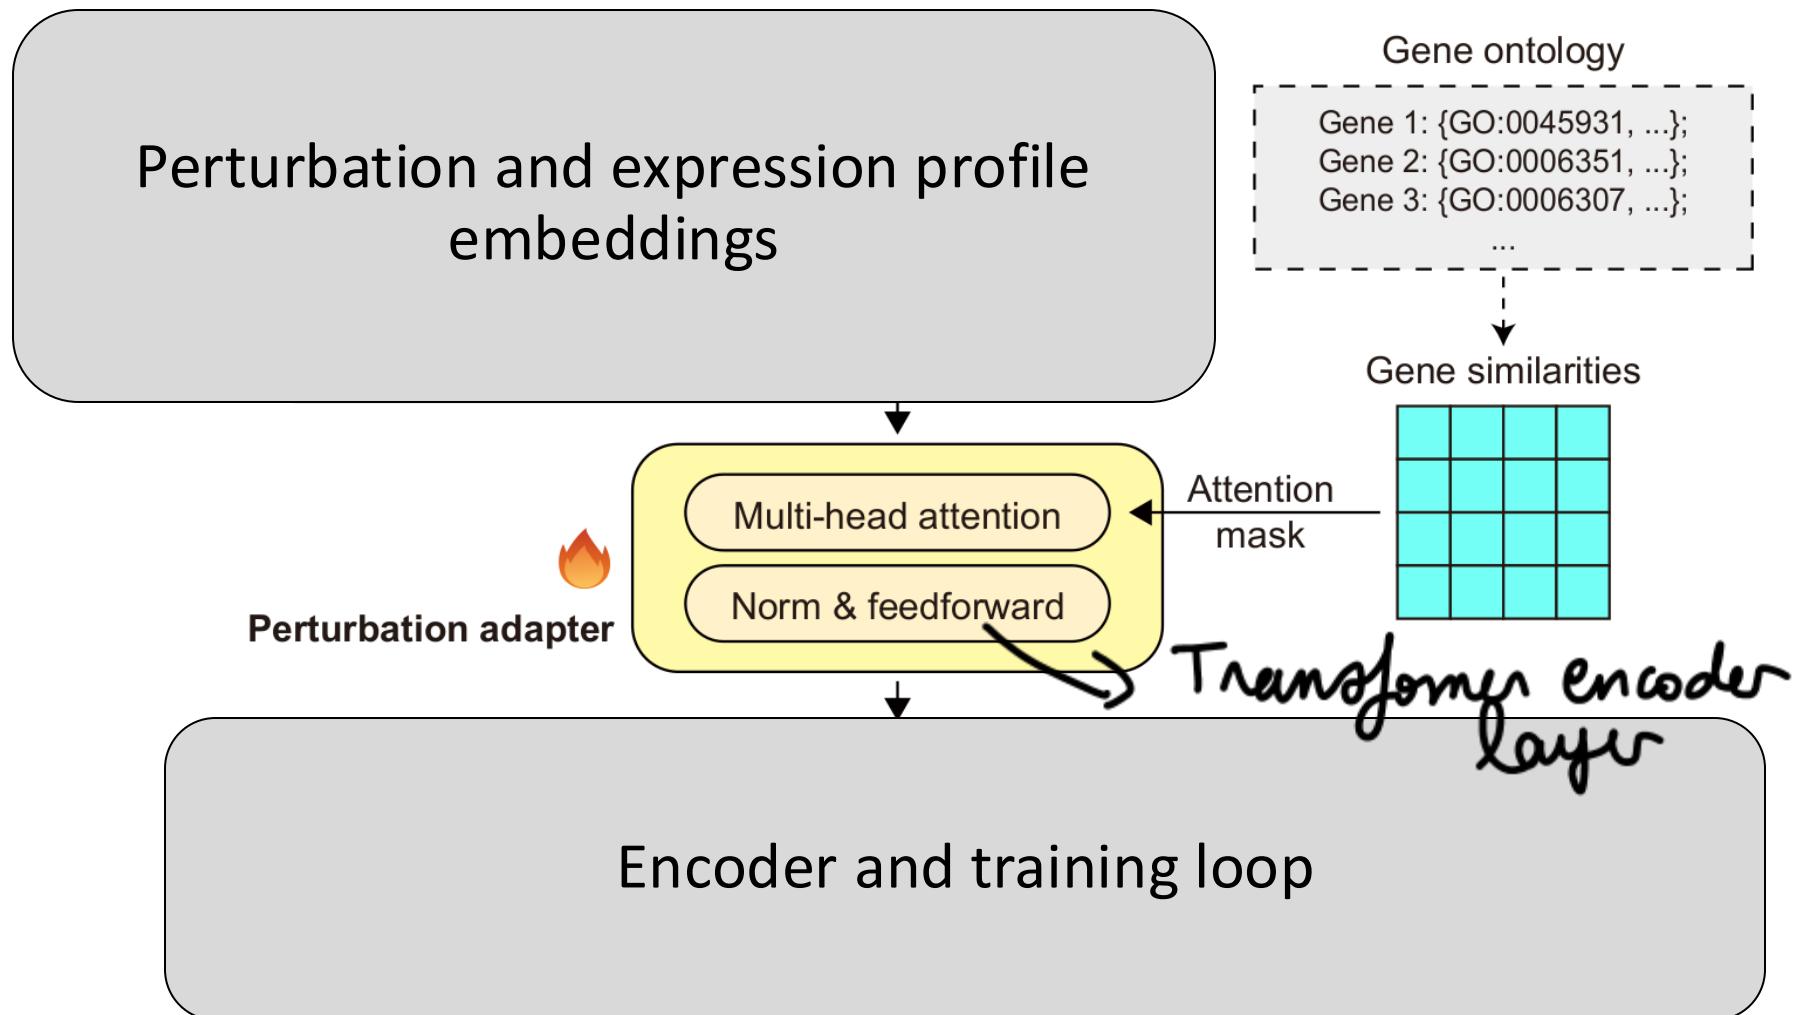
Plug-in perturbation adapter downstream-task tuning and biological prior integration



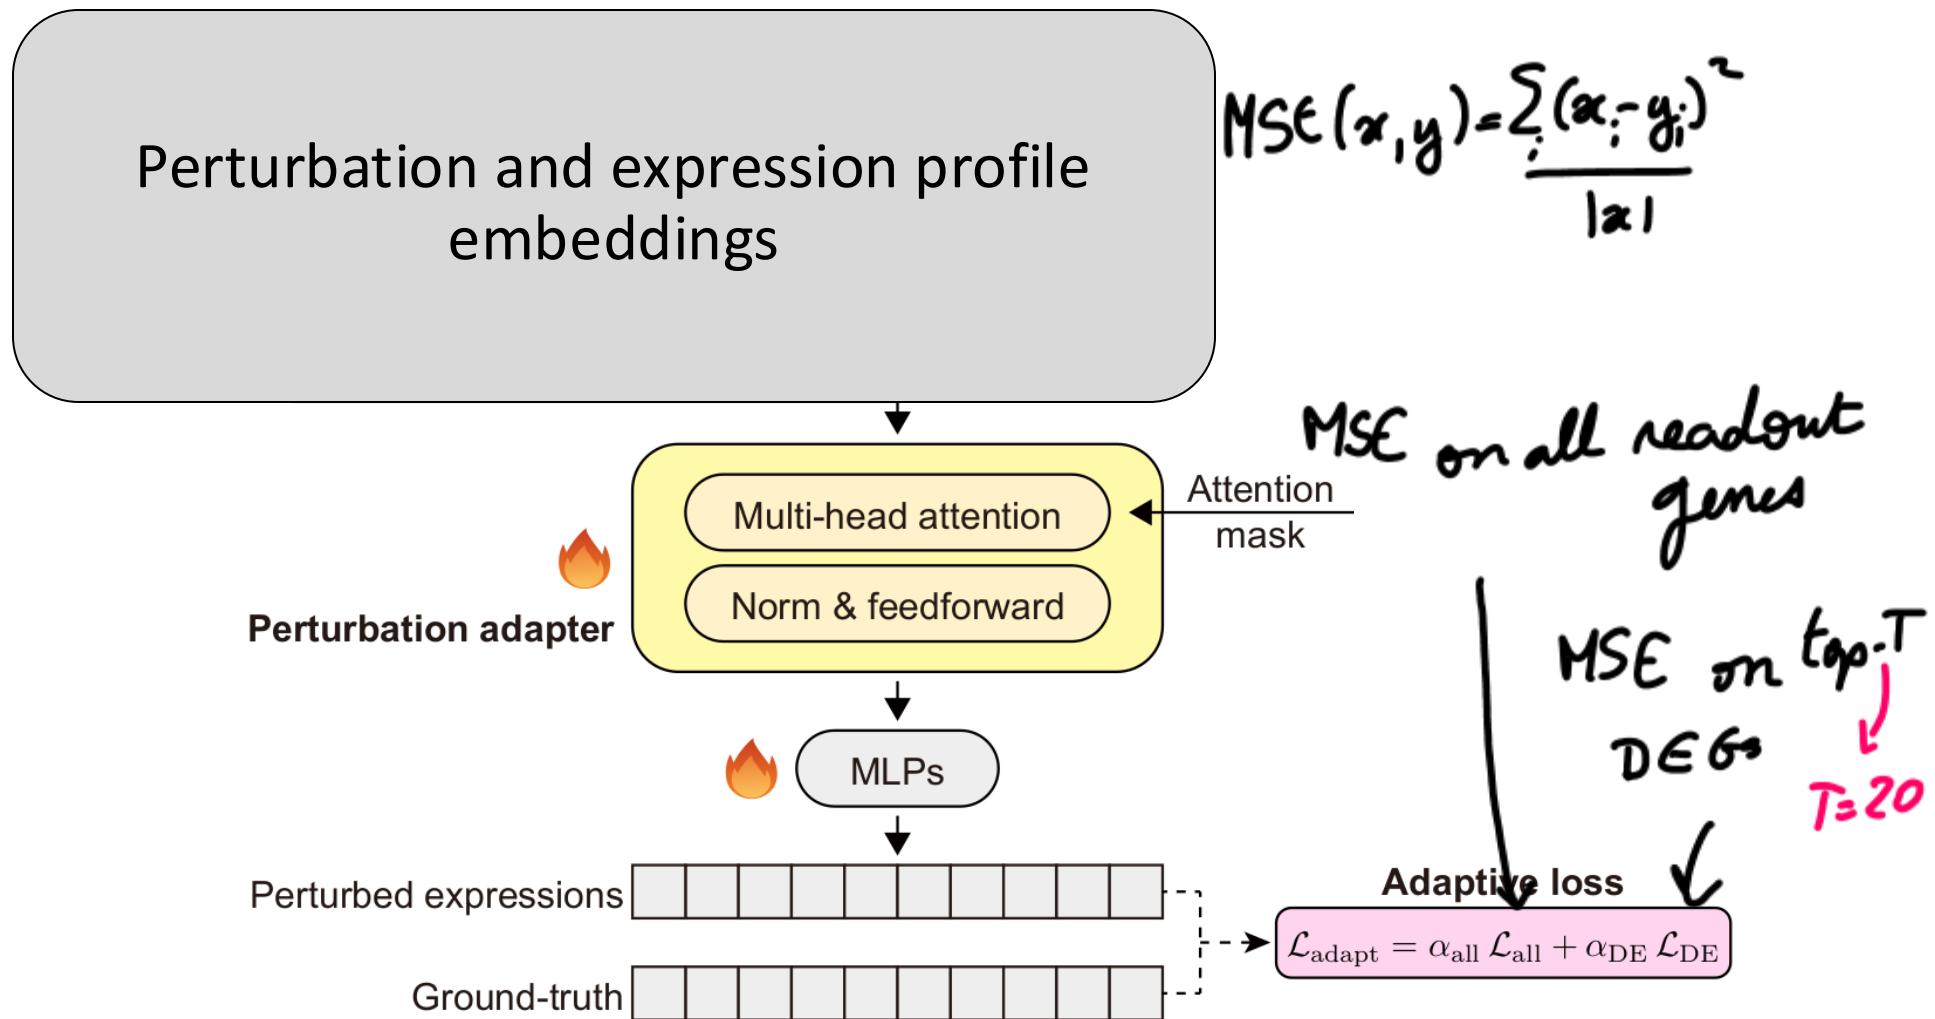
Plug-in perturbation adapter downstream-task tuning and biological prior integration



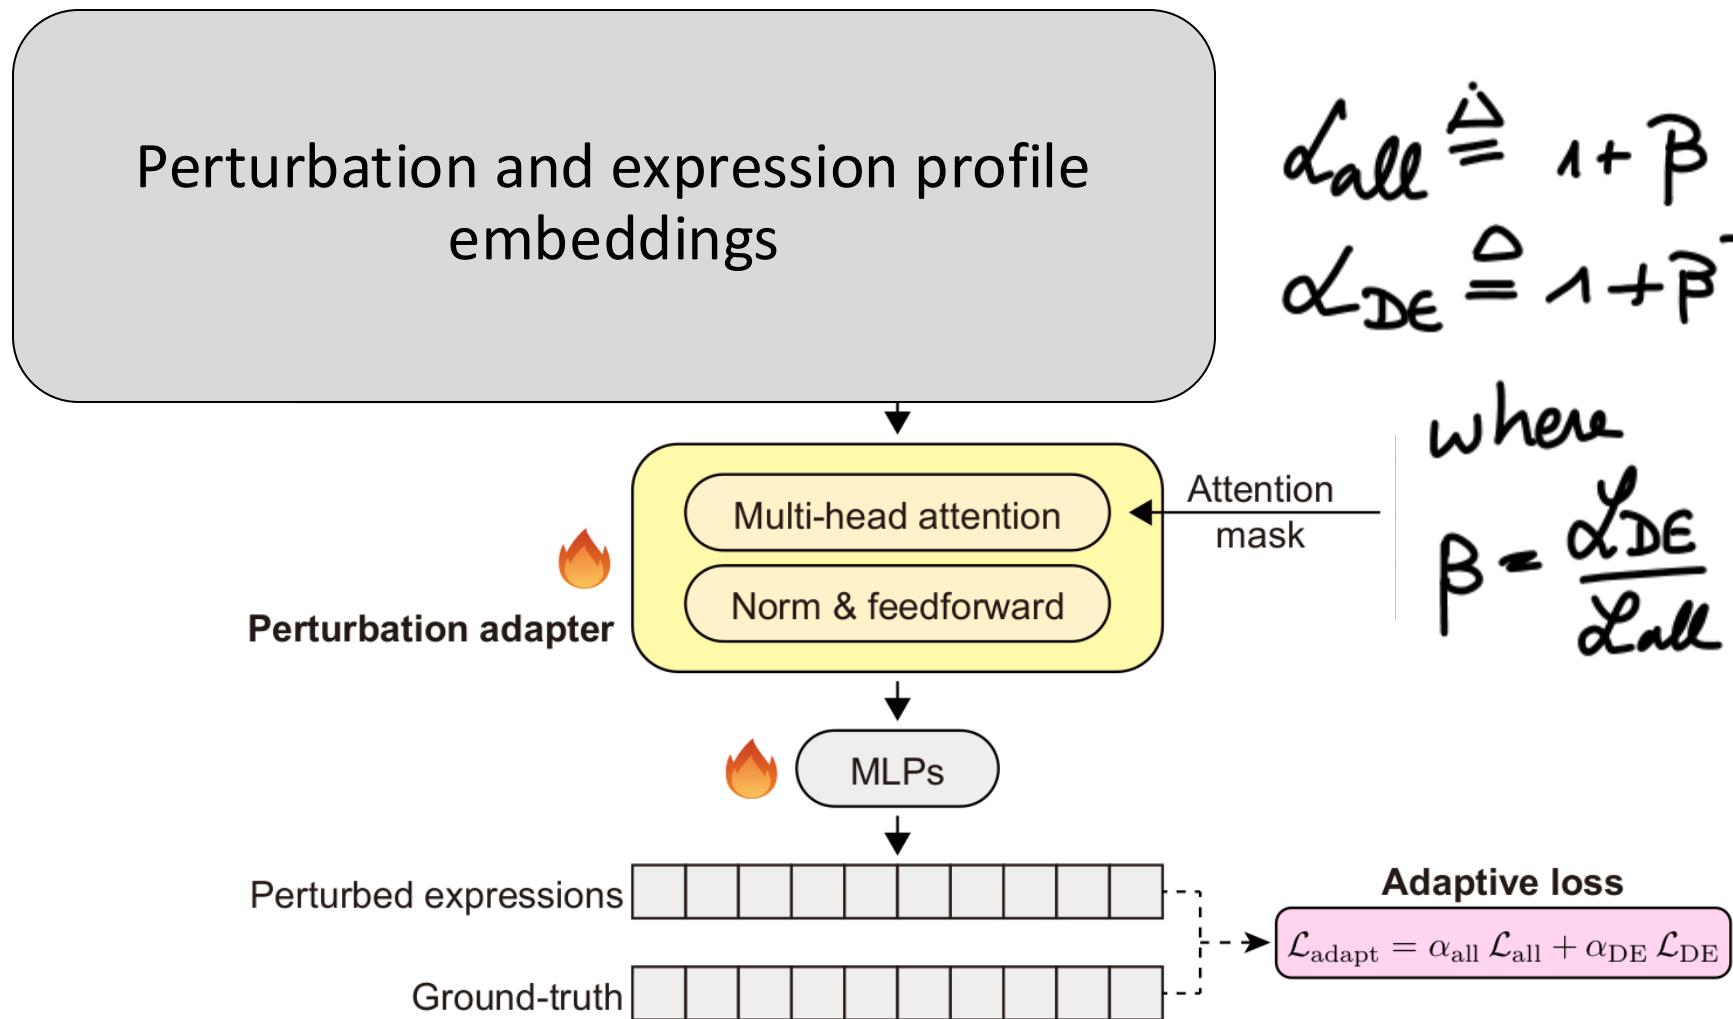
Plug-in perturbation adapter downstream-task tuning and biological prior integration



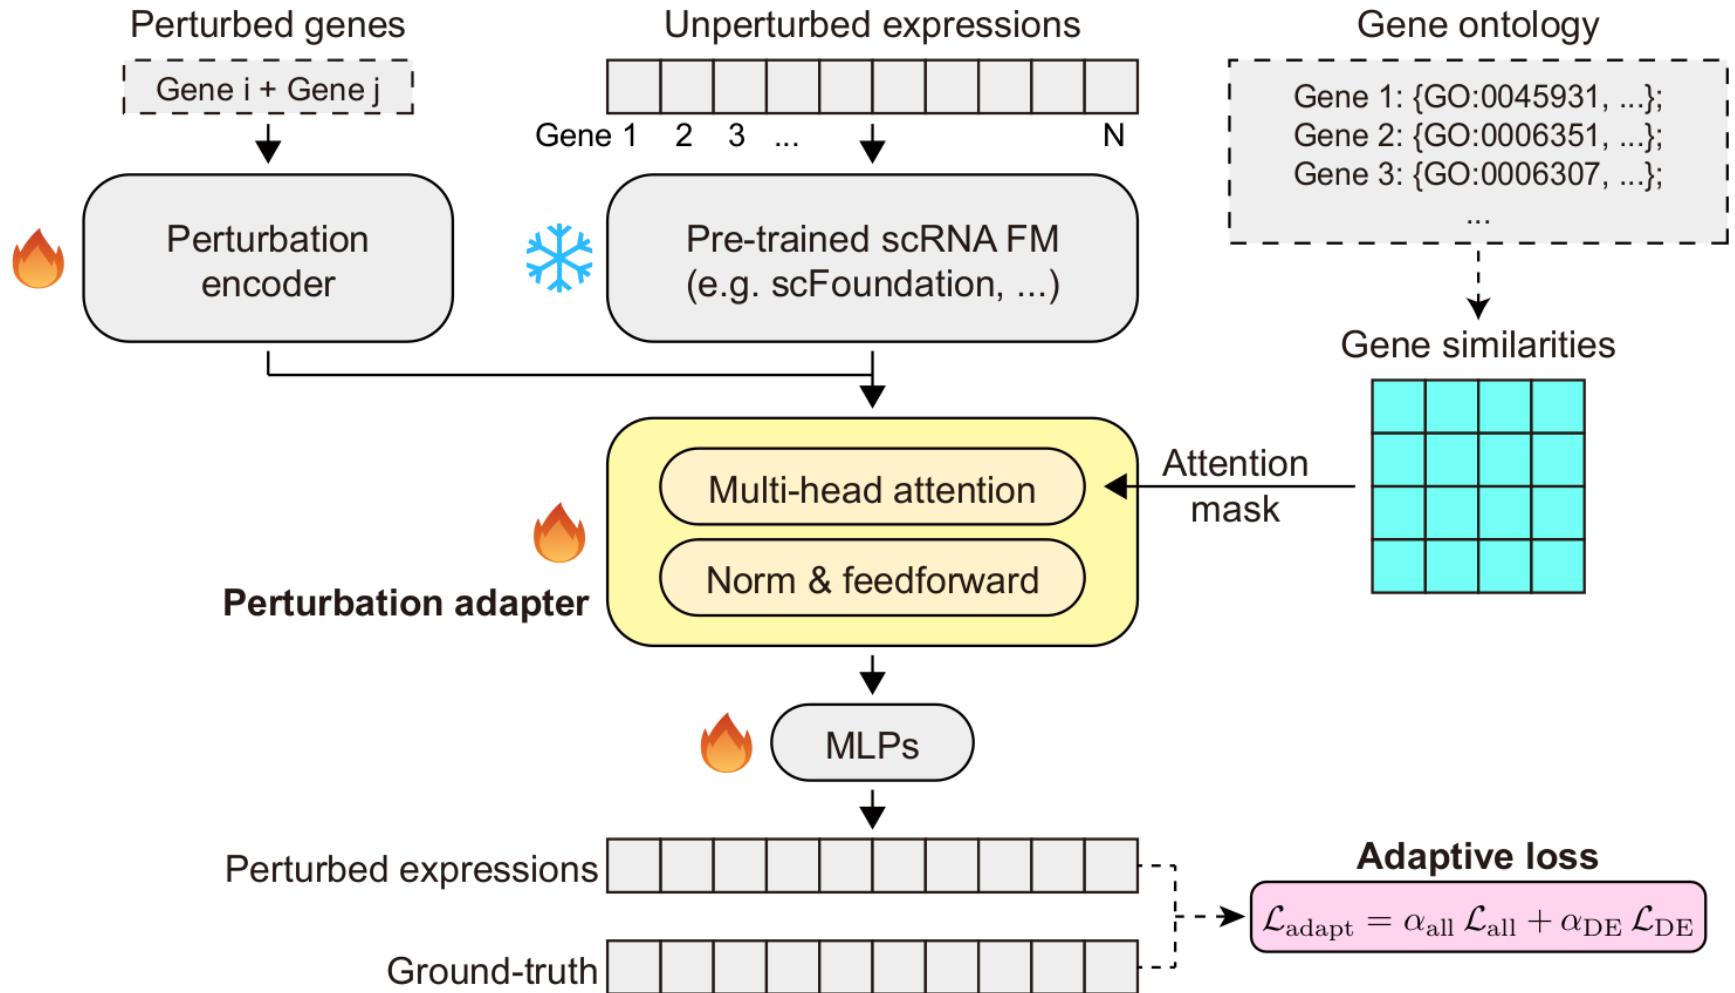
Adaptive loss put more weight on differentially expressed genes to lessen noise



Plug-in perturbation adapter downstream-task tuning and biological prior integration



Final architecture of PertAdapt



Experiments: Prediction of post-perturbation expression showing only single-gene perturbations

Model	Reprogle RPE1		Reprogle K562		Nadig HepG2	
	MSE↓	Pearson↑	MSE↓	Pearson↑	MSE↓	Pearson↑
No-Change	0.201	0.706	0.193	0.908	0.121	0.681
Linear Model	0.113	0.704	0.165	0.924	0.101	0.651
GEARS	0.159	0.719	0.124	0.929	0.101	0.652
scGPT	0.158	0.721	0.135	0.919	0.110	0.676
GeneCompass	0.170	0.727	0.121	0.921	0.104	0.680
scFoundation	0.166	<u>0.731</u>	0.142	0.927	0.109	0.684
AIDO.Cell	0.134	0.724	0.130	0.928	0.102	0.680
scFoundation+Ours	0.091	0.737	0.083	<u>0.935</u>	0.062	0.700
AIDO.Cell+Ours	<u>0.094</u>	0.725	<u>0.088</u>	0.936	<u>0.067</u>	<u>0.685</u>
	Nadig Jurkat		Adamson		Dixit	
	MSE↓	Pearson↑	MSE↓	Pearson↑	MSE↓	Pearson↑
No-Change	0.111	0.801	0.905	0.787	0.380	0.831
Linear Model	0.101	0.815	0.675	0.862	0.229	0.912
GEARS	0.095	0.817	1.028	0.795	0.271	0.896
scGPT	0.112	0.806	1.116	0.679	0.295	0.841
GeneCompass	0.099	0.808	0.986	0.808	0.257	0.890
scFoundation	0.131	0.812	0.792	0.806	0.243	0.913
AIDO.Cell	0.103	0.820	0.804	0.840	0.213	0.924
scFoundation+Ours	0.061	0.826	0.324	0.904	<u>0.021</u>	<u>0.992</u>
AIDO.Cell+Ours	<u>0.069</u>	<u>0.824</u>	0.408	<u>0.886</u>	0.013	0.995

Experiments: Prediction of post-perturbation expression showing only single-gene perturbations

Model	Reprogle RPE1		Reprogle K562		Nadig HepG2		
	MSE↓	Pearson↑	MSE↓	Pearson↑	MSE↓	Pearson↑	
No-Change	0.201	0.706	0.193	0.908	0.121	0.681	
Linear Model	0.113	0.704	0.165	0.924	0.101	0.651	
GEARS	0.159	0.719	0.124	0.929	0.101	0.652	
scGPT	0.158	0.721	0.135	0.919	0.110	0.676	
GeneCompass	0.170	0.727	0.121	0.921	0.104	0.680	
scFoundation	0.166	<u>0.731</u>	0.142	0.927	0.109	0.684	
AIDO.Cell	0.134	0.724	0.130	0.928	0.102	0.680	
scFoundation+Ours	0.091	0.737	0.083	<u>0.935</u>	0.062	0.700	
AIDO.Cell+Ours	<u>0.094</u>	0.725	<u>0.088</u>	0.936	<u>0.067</u>	<u>0.685</u>	
		Nadig Jurkat	Adamson		Dixit		
		MSE↓	Pearson↑	MSE↓	Pearson↑	MSE↓	Pearson↑
No-Change	0.111	0.801	0.905	0.787	0.380	0.831	
Linear Model	0.101	0.815	0.675	0.862	0.229	0.912	
GEARS	0.095	0.817	1.028	0.795	0.271	0.896	
scGPT	0.112	0.806	1.116	0.679	0.295	0.841	
GeneCompass	0.099	0.808	0.986	0.808	0.257	0.890	
scFoundation	0.131	0.812	0.792	0.806	0.243	0.913	
AIDO.Cell	0.103	0.820	0.804	0.840	0.213	0.924	
scFoundation+Ours	0.061	0.826	0.324	0.904	<u>0.021</u>	<u>0.992</u>	
AIDO.Cell+Ours	<u>0.069</u>	<u>0.824</u>	0.408	<u>0.886</u>	0.013	0.995	

Experiments: Ablation study adapted from the SimCLR framework

Backbone	<i>adapter attention loss</i>				Norman (Overall)		Replogle RPE1	
	FM	PA	AM	\mathcal{L}_{adapt}	MSE↓	Pearson↑	MSE↓	Pearson↑
scFoundation	✓				0.375	0.584	0.328	0.490
					0.210	0.776	0.166	0.731
	✓	✓			0.208	0.779	0.143	0.730
	✓	✓	✓		0.183	<u>0.782</u>	0.124	<u>0.733</u>
	✓			✓	<u>0.178</u>	<u>0.782</u>	<u>0.109</u>	0.732
AIDO.Cell	✓	✓	✓	✓	0.156	0.789	0.091	0.737
					0.351	0.618	0.297	0.515
	✓				0.192	0.781	0.134	0.724
	✓	✓			0.197	0.769	0.125	0.720
	✓	✓	✓		0.179	<u>0.782</u>	0.104	0.723
	✓			✓	<u>0.170</u>	0.780	<u>0.103</u>	0.726
	✓	✓	✓	✓	0.159	0.786	0.094	<u>0.725</u>

Perspectives

1. Comments on the paper
2. Why is it interesting for BioComp?

My comments on the paper

Strengths:

- Convincing results
- I liked the plug-in perturbation adapter and GO-based attention mask
- Open-source: <https://github.com/BaiDing1234/PertAdapt>

Weaknesses:

- Lack of biological assessment
- Lack of runtime assessment
- "Biological prior vs exploration" dilemma

Your comments?

Why is it interesting for BioComp?

Multiple applications (e.g., drug discovery)
Many of you working on this topic

- Lots of papers on the topic
- Often not so much biology-focused
- Interpretability, practicality (time+resources),
probably other challenges